

CEO Local AI Platform

An always-on personal AI workstation designed to run locally on Apple Silicon, remove hosted token limits, and combine strong reasoning, tools, memory, and automation in one system.

Important: this page is a product vision. The current repo implements a Gemini-based CEO assistant, not a finished local LLM stack.

Who It Is For

A power user, founder, or developer who wants one private AI system that is always available, runs from a desk, and acts more like an operator than a chatbot.

Why This Exists

- No waiting for hosted token resets or usage caps.
- Local uptime on a Mac mini or MacBook Pro that can stay online 24/7.
- Stronger continuity through persistent memory, tools, and personal context.
- One interface for coding, planning, notes, GitHub, and device actions.

Current Starting Point

Use the existing CEO app as the shell: mobile client, FastAPI backend, WebSocket transport, voice I/O, and tool integrations. Swap the hosted model layer for a local runtime over time.

Target Capabilities

- Fast local inference with a primary model tuned for desktop use.
- A planner/executor layer that can route tasks across coding, research, and automation flows.
- Native tool use for files, git, notes, browser tasks, shell commands, and APIs.
- Long-lived memory across sessions, projects, and personal preferences.
- Always-on voice plus text interaction from phone or desktop.
- Built-in evals to score outputs against speed, accuracy, and task completion.
- Optional cloud escalation only when local capability is insufficient.

Compact Architecture

1. Client layer: mobile app, desktop UI, voice input/output.
2. Local gateway: session manager, prompt assembly, auth, and observability.
3. Model runtime: local inference server on Apple Silicon with quantized models.
4. Orchestration: planner, tool router, memory retrieval, and task queue.
5. Data plane: notes, repos, embeddings, logs, and evaluation results.

How It Would Run

1. Host the backend and model runtime on a Mac mini or MacBook Pro with Apple Silicon.
2. Keep CEO running as a local service so the phone app connects over LAN by default.
3. Store memory, logs, and project state locally first; sync selectively when needed.
4. Use cloud models only as fallback, not as the default execution path.